

Teaching Monocular Depth the Metre with the Phone’s Own LiDAR: Fine-Tuned vs. Pretrained Depth Anything V2 for Room Reconstruction, Measured Against a Laser Scanner

Author Name
Affiliation
email@example.com

Abstract—Devices without LiDAR can reconstruct a room from monocular depth, but an off-the-shelf “metric” model such as Depth Anything V2 Metric-Indoor, trained on synthetic images, misplaces a real room by 54.3 cm inside a reconstruction pipeline. We ask whether fine-tuning it with depth from the phone’s own LiDAR—no laser scanner involved—makes it better than its pretrained self when measured against the Faro laser scanner that ARKitScenes ships. We fine-tune only the DPT head on 24 ARKitScenes rooms with iPad LiDAR labels (R3) and compare it with the pretrained checkpoint through one depth $\rightarrow$ TSDF $\rightarrow$ mesh pipeline in which the depth source is the single variable, on rooms never used for training: six rooms with a Faro reference mesh (3D and 2D) and 20 further rooms (2D, 2,090 frames). R3 lowers Chamfer from 54.3 cm to 15.7 cm (3.5 $\times$, better in every room) and per-frame AbsRel from 0.410 to 0.130 on the 20 unseen rooms (better in 19/20), with no calibration or sensor at run time. A laser teacher (14.6 cm) is not measurably better. For context, the same pipeline gives 3.0 cm with iPad LiDAR and 6.0 cm with a monocular model scaled per frame on $\sim$ 200 sparse metric points. Per-frame analysis shows that R3’s remaining error is a scale that swings with the content of each view, which fine-tuning de-biases but does not remove: the rooms are a large step toward, but not yet at, overview accuracy. Code, configs, per-run metrics and checkpoints are released.

Index Terms—monocular depth estimation, transfer learning, metric depth, LiDAR supervision, TSDF fusion, 3D room reconstruction, ARKitScenes

I. INTRODUCTION AND BUSINESS MOTIVATION

Scanning a room into a 3D model is already a consumer feature on iPhone Pro and iPad Pro (Apple RoomPlan, Polycam), all relying on the *LiDAR sensor* for metric depth. Devices with LiDAR are, however, a minority: non-Pro iPhones and almost every Android phone have none. Their alternative is *monocular depth estimation*, which advanced rapidly between 2020 and 2024 (MiDaS, Depth Anything). A product team must know **by how many centimetres the room is wrong if LiDAR is replaced by a monocular model, and which use cases that still serves**—a question per-image benchmark metrics (AbsRel, δ_1) do not answer, and end-to-end reconstruction systems (Atlas, NeuralRecon, SimpleRecon) cannot answer

because they design the whole pipeline around their own model.

Some recent models claim to output metres directly (DA-v2 Metric-Indoor, ZoeDepth, Depth Pro), but they are trained on data that is not a phone’s camera—DA-v2 Metric-Indoor on rendered Hypersim images—and inside a real reconstruction pipeline the room comes out 54.3 cm wrong (§V-C). The direct remedy is *transfer learning* on target-domain depth. The question is where the labels come from: a \$20k laser scanner, or the *LiDAR already inside Pro phones*, which an app maker can collect at scale. If the latter works, devices with LiDAR can teach those without it.

Research question. Is DA-v2 Metric-Indoor fine-tuned with iPad LiDAR depth (no laser scanner at training time) better than the same model pretrained, measured against the Faro laser scanner—as a room mesh and as per-frame depth—on rooms never used for training? **Supporting questions:** (a) is a laser teacher better than a phone-LiDAR teacher; (b) where does the fine-tuned model stand relative to a sensor at run time; (c) where does the remaining error come from?

Approach. We build a classical pipeline (per-frame depth $\rightarrow$ scale alignment $\rightarrow$ TSDF fusion $\rightarrow$ mesh) in which **the depth source is the single swappable variable**, so “pretrained” and “fine-tuned” are two rows differing only in a checkpoint, and every other depth source (Faro laser, ARKit LiDAR, relative monocular models with different scale strategies) is priced through the same pipeline as context.

Contributions.

- 1) A measured answer to whether fine-tuning a metric monocular model on a phone’s own LiDAR helps, against a laser scanner on held-out rooms in 3D (6 rooms) and 2D (6 + 20 rooms), with checkpoints and room split released.
- 2) A teacher ablation (Faro vs. LiDAR, 10 vs. 24 rooms) and an intrinsics-conditioned baseline (Depth Pro).
- 3) An open-source pipeline in which DepthSource and ScaleAligner are interfaces switched from config, so every table is “one loop, one variable.”

- 4) “Centimetres on the wall” for laser / LiDAR / monocular $\times$ scale strategy, and an analysis showing that monocular error comes from per-view scale, not temporal drift—explaining why fine-tuning helps and why it is not enough.

Not claimed. We do not compare against Atlas / NeuralRecon / SimpleRecon; we do not claim fine-tuning makes LiDAR unnecessary (R3 at 15.7 cm remains far from LiDAR at 3.0 cm); and since training and testing use the same iPad camera, transfer to other phone cameras is not measured.

II. RELATED WORK

Relative monocular depth. MiDaS [1] trains across datasets with a scale- and shift-invariant loss, so its output is disparity up to an affine transform: two constants (s, t) per image are needed for metres. Depth Anything [2] and V2 [3] scale this idea with large unlabelled corpora. We treat (s, t) as an experimental variable (the `ScaleAligner`) rather than hiding it in the model.

Metric monocular depth. ZoeDepth [4] and the metric DA-v2 variants (fine-tuned on Hypersim / Virtual KITTI) output metres directly. Metric3D [5], UniDepth [6] and Depth Pro [7] instead condition on, or estimate, camera intrinsics. These works report per-image metrics on NYUv2/KITTI; we show that on real room video the per-frame scale of a metric model still swings widely within one scan—invisible to per-image metrics, obvious once frames are fused. The closest work in data is Prompt Depth Anything [8], which feeds iPhone LiDAR to the network *at run time* on ARKitScenes; we instead use LiDAR only as a *training* teacher, so the deployed model needs no depth sensor.

Depth fusion and learned reconstruction. TSDF fusion [9], [10] neither knows nor cares where metric, posed depth comes from, making it a fair fixed stage for every source; we use Open3D [11]. Atlas [12], NeuralRecon [13] and SimpleRecon [14] learn depth or TSDF from multiple images and are strong ScanNet baselines, but do not answer what room a per-image model plus the device’s VIO pose yields, so we do not compare against them. Scale recovery from sparse tracked points follows CNN-SLAM [15] and MonoSDF [16]; our `sparse_points` aligner belongs to this line.

Mobile LiDAR and products. ARKitScenes [17] is the first dataset releasing both iPad LiDAR depth and Faro laser depth for the same rooms. RoomPlan [18] and similar apps require LiDAR, and no public report states how many centimetres they deviate from a laser scanner.

III. SYSTEM DESIGN

The research question (“change only the depth source—how much changes?”) leads to one design rule: every depth source passes through the same code after the point where it is produced, and whatever can be swapped is swapped from a YAML config, so one result table is one config file (Fig. 1).

DepthSource. An interface `get_depth(frame)` returning an (H, W) depth map and a flag `is_metric`. Implementations read Faro depth (`gt`), ARKit LiDAR (`lidar`) or

TABLE I
SCALE ALIGNERS. EVERY NON-IDENTITY ALIGNER FITS (s, t) IN INVERSE DEPTH WITH TRIMMED LEAST SQUARES.

Aligner	Fitted on / meaning
<code>identity</code>	nothing; for Faro, LiDAR and metric models
<code>oracle_frame</code>	that frame’s GT; the model’s ceiling (not deployable)
<code>per_scene</code>	GT of 10 frames spread over the scan, fitted once; upper bound of a one-off calibration
<code>sparse_points</code>	200 LiDAR pixels per frame at confidence 2, a proxy for VIO feature points; upper bound of a VIO-based system

run a monocular network (`mono`: DA-v2, MiDaS, fine-tuned checkpoints) on the upright RGB frame.

ScaleAligner. A relative model is correct only up to an affine map, and *how* that map is obtained is a system decision measured separately from the model (Table I). Affine-invariant models are affine in *disparity*, so every aligner fits $1/gt \approx s \cdot (1/pred) + t$ by trimmed least squares and inverts back; on the development scene this change alone moved `mono+oracle` from 15.2 cm to 5.4 cm. Metric sources must use `identity`; the pipeline refuses to build otherwise.

Fixed stages. Open3D `ScalableTSDFVolume` with 4 cm voxels, truncation $3 \times$ voxel and depth range 0.1–5 m; every source is resampled to the LiDAR resolution (256×192) for fairness, and only pixels where Faro has depth are integrated, so no source is penalised for seeing beyond the reference. Marching cubes and small-cluster removal follow. Poses are the dataset’s ARKit VIO trajectory for every row.

IV. EXPERIMENTAL SETUP

A. Dataset and reference

ARKitScenes [17] is the only dataset that provides, for the same room, RGB video from a real device (iPad Pro, 640×480), that device’s LiDAR depth (256×192) and Faro laser depth registered to the camera poses (1920×1440 , ~ 4 fps). **Ground truth is the Faro data as distributed.** In 2D we compare per-frame depth with Faro depth directly (nearest-neighbour down-sampling, so every GT value is a measured one). In 3D, since no laser mesh is shipped, we fuse Faro depth from every frame at 1 cm voxels once per scene into a reference mesh. The `gt` row at 4 cm therefore has non-zero error—the pipeline’s own discretisation loss—and because the reference uses the same poses as every row, pose drift is excluded and all differences are due to the depth source. iPad LiDAR is never used as ground truth.

The six test scenes (Table II) span standard, large and glossy/mirror rooms. Scene 42444474 was used during development, so its numbers are not held-out; per-scene results show the conclusions hold on the other five.

B. Fine-tuning protocol

Scenes disjoint from the test set at the `visit_id` level were drawn with seed 0: 10 training scenes with both Faro




Fig. 1. Pipeline. The two highlighted stages are the only things that change between rows of a result table; everything to their right is fixed.

TABLE II
TEST SCENES (3D AND 2D). FRAMES = FARO FRAMES EVERY ROW IS EVALUATED ON.

video_id	Room	Extent (m)	Frames
42444474	kitchen/dining, glass (dev. scene)	$9.9 \times 5.1 \times 4.2$	236
47333774	standard bedroom	$5.1 \times 4.9 \times 3.3$	375
47115299	large living room	$8.1 \times 6.1 \times 2.5$	324
42897743	bathroom (mirror, tiles)	$8.0 \times 3.4 \times 3.6$	448
47429736	bedroom + closets	$6.6 \times 6.4 \times 2.7$	504
45261556	kitchen (glossy cabinets)	$7.1 \times 6.1 \times 2.4$	465

and LiDAR, 14 more with LiDAR only, and 3 Validation-fold scenes for checkpoint selection. Three fine-tunes share one recipe: **R1** uses Faro labels on the 10 scenes (3594 frames), **R2** iPad LiDAR (confidence ≥ 1) on the same 10 scenes (8656 frames), and **R3, the main model**, LiDAR on all 24 scenes (19 896 frames). R1 and R2 do not share frames: Faro exists at ~ 4 fps, LiDAR with every colour frame (we take every third, ~ 10 fps), so R2 sees $\sim 2.4\times$ more distinct views than R1. The DINOv2 encoder is frozen; only the DPT neck and metric head train, with SiLog loss, AdamW (lr 5×10^{-5} , cosine schedule), batch 8 (2×4 accumulation) on 518×518 crops, for 6000 steps (R1, R2) or 12 000 (R3), keeping the checkpoint with the lowest validation AbsRel. A 6000-step run takes about an hour on an RTX 3070. At test time the fine-tuned models are plain DepthSource rows: RGB in, metres out, identity aligner, same input preparation as the pretrained model.

C. Metrics and experiments

3D: 200k points sampled on both meshes; accuracy = mean pred $\rightarrow$ ref distance, completeness = ref $\rightarrow$ pred, Chamfer = their mean; precision/recall/F-score at $\tau = 2/5/10$ cm, with F@5 cm as headline (5 cm $\approx$ furniture-placement tolerance). **2D:** per-frame AbsRel, δ_1 and RMSE against Faro depth; pixels are selected by the ground truth (0.1–5 m) and predictions clipped to that range, averaged per room then over rooms. Table III lists the experiments; numbers are mean $\pm$ std over the six scenes. **Exp. 7** adds 20 Validation-fold rooms (seed 0, one video per visit, excluding every visit of the training, validation and test sets) evaluated in 2D only on every second Faro frame; these rooms decided nothing.

TABLE III
EXPERIMENTS. EXP. 6–7 ANSWER THE MAIN QUESTION; EXP. 1–5 GIVE CONTEXT.

Exp.	Variable (everything else fixed)
6	pretrained vs. R3 ; ablation R1 / R2 / Depth Pro (6 scenes, 3D + 2D)
7	pretrained / R3 / LiDAR on 20 further rooms (2D)
1	depth source $\times$ aligner: gt / lidar / DA-v2 L + {oracle, sparse, per_scene} / metric
2–5	voxel size; frame stride; model size; LiDAR confidence




Fig. 2. Per-room AbsRel against Faro: pretrained (x) vs. R3 (y), six test rooms and 20 Validation-fold rooms. Below the diagonal = the fine-tune is better.

V. RESULTS

All numbers are generated from the per-run metrics.json files; units are cm and F is F@5 cm.

A. Main result: R3 vs. the pretrained model

The two rows differ only in the weights of the DPT neck and metric head; at test time both see the same frames and input with no alignment of any kind, and no room here was used for training or checkpoint selection.

TABLE IV

MAIN RESULT ON THE SIX TEST ROOMS: RGB ONLY, NO ALIGNMENT, MEASURED AGAINST FARO. THE TWO REFERENCE ROWS (EXP. 1) NEED A SENSOR AT RUN TIME.

Model	Teacher	Chamfer	Acc.	Compl.	F@5	R@10	AbsRel	δ_1	RMSE
DA-v2 Metric-Indoor, pretrained	— (Hypersim)	54.3 $\pm$ 10.6	70.3	38.3	0.03	0.10	0.382	0.277	49.3
+ fine-tune R3	iPad LiDAR, 24 rooms	15.7 $\pm$ 4.0	17.3	14.0	0.24	0.51	0.134	0.812	20.2
ref.: iPad LiDAR (sensor)	—	3.0 $\pm$ 0.6	2.2	3.9	0.93	0.97	0.018	0.997	3.9
ref.: DA-v2 L + 200 sparse pts/frame	—	6.0 $\pm$ 1.7	6.4	5.5	0.74	0.91	0.036	0.985	7.4

TABLE V
EXP. 6 PER TEST ROOM (ORDER OF TABLE II): CHAMFER (CM) AND ABSREL.

	4244	4733	4711	4289	4742	4526	mean
Depth Pro, Chamfer	131.1	165.2	172.9	196.4	177.9	88.9	155.4
pretrained, Chamfer	48.5	61.6	41.5	48.9	54.1	71.0	54.3
R1 (Faro), Chamfer	10.8	13.8	17.6	21.1	12.4	11.9	14.6
R2 (LiDAR), Chamfer	13.9	19.0	12.8	12.5	16.6	24.6	16.6
R3 (LiDAR$\times$24), Chamfer	12.8	16.7	17.3	12.2	12.4	22.5	15.7
pretrained, AbsRel	0.375	0.346	0.255	0.412	0.476	0.430	0.382
R3, AbsRel	0.127	0.156	0.193	0.087	0.141	0.101	0.134

TABLE VI
EXP. 7: 2D ONLY, 20 VALIDATION-FOLD ROOMS NEVER USED FOR ANYTHING ELSE (2,090 FRAMES); MEAN $\pm$ STD OVER ROOMS.

Row	AbsRel	δ_1	RMSE (cm)
DA-v2 Metric-Indoor, pretrained	0.410 $\pm$ 0.164	0.333 $\pm$ 0.216	50.9
+ fine-tune R3	0.130 $\pm$ 0.046	0.827 $\pm$ 0.135	18.4
ref.: iPad LiDAR (sensor)	0.019 $\pm$ 0.002	0.997 $\pm$ 0.002	3.8

- 1) **R3 beats the pretrained model in every test room.** Chamfer 54.3 cm $\rightarrow$ 15.7 cm (3.5 $\times$; 38.6 $\pm$ 8.6 cm better per room, 6/6 rooms, exact Wilcoxon $p = 0.031$, the smallest attainable at $n = 6$) and F@5cm 0.03 $\rightarrow$ 0.24 (Tables IV, V). Accuracy falls more than completeness: the pretrained model builds the room too far away (15–31 % over-estimation, §VI-B); R3 brings surfaces back near where they belong.
- 2) **The same holds on rooms that took part in nothing.** On the 20 Validation-fold rooms AbsRel drops 0.410 $\rightarrow$ 0.130 (3.2 $\times$; better in 19/20 rooms, exact Wilcoxon $p = 3.8 \times 10^{-6}$) and δ_1 rises 0.333 $\rightarrow$ 0.827 (Table VI, Fig. 2). The one exception, 48018730 (AbsRel 0.275 $\rightarrow$ 0.281), is one of the pretrained model’s best rooms; in the other 19, R3’s AbsRel is 21–64 % of the pretrained one.
- 3) **It does not reach a sensor.** R3 is $\sim$ 13 cm behind iPad LiDAR in Chamfer and 9.7 cm behind a per-frame fit on 200 sparse points; fine-tuning fixes the scale of the domain, not of the frame being looked at (§VI-A).

B. Teacher ablation and a no-fine-tuning baseline

The laser teacher is not measurably better than the LiDAR teacher. R1 14.6 cm, R3 15.7 cm and R2 16.6 cm all lie within the across-room std (4.0–4.7 cm), and no row wins every room (R1 beats R3 in three rooms, R3 wins two, one tie; Table V, Fig. 3). Because R1 and R2 also differ in the number




Fig. 3. Exp. 6 top-down error maps (rows = scenes; columns = Depth Pro, the three fine-tunes and the pretrained metric model, as labelled above each panel; colour = distance to the reference, 0–10 cm). The pretrained column is red almost everywhere—the room is built at the wrong distance; after fine-tuning the error concentrates at object edges and far walls.

of views, the honest claim is that a phone’s depth sensor can replace the laser as the teacher at no measurable cost on these rooms—not that the two are equal. Fourteen extra LiDAR-only rooms (R3 vs. R2) do not help measurably either; we keep R3 as the main model because it was best on validation and is the only recipe that scales without a laser.

Knowing the camera does not rescue metric depth. Depth Pro [7], given the true focal length and no fine-tuning, reaches 155.4 cm (F@5 cm 0.00), worse than the pretrained DA-v2 Metric row: it predicts a 48° field of view where the camera has 62°, and even after removing all scale per frame its AbsRel is 0.24–0.54. We verified the wrapper on a synthetic room with exact ground truth; the claim is about this model on this

TABLE VII
EXP. 1: SAME PIPELINE, ONLY `DEPTH.SOURCE` / `DEPTH.ALIGNER`
CHANGE (6 SCENES).

Source	Aligner	Chamfer	F@2	F@5	AbsRel
Faro (gt)	identity	2.1 ± 0.5	0.86	0.97	0
ARKit LiDAR	identity	3.0 ± 0.6	0.54	0.93	0.018
DA-v2 L	oracle_frame	5.3 ± 1.3	0.46	0.79	0.032
DA-v2 L	sparse_points	6.0 ± 1.7	0.38	0.74	0.036
DA-v2 L	per_scene	22.9 ± 2.1	0.08	0.23	0.237
DA-v2 Metric	identity	54.3 ± 10.6	0.01	0.03	0.382

TABLE VIII
SECONDARY EXPERIMENTS (CHAMFER IN CM, 6 SCENES).

Exp.	Result
2 voxel (gt)	2 / 4 / 8 cm: 1.2 / 2.1 / 2.8; 4 cm is used throughout (fusion 1–3 s per scan at every size)
3 frame stride	GT accuracy is constant (1.2–1.3 cm); fewer frames only lose coverage: F@5 cm 0.97 at 4 fps, 0.91 at 0.8 fps, 0.83 at 0.4 fps
4 model size (oracle)	DA-v2 L / S / MiDaS small: 5.33 / 6.35 / 13.24 at 871 / 121 / 27 ms per frame
5 LiDAR confidence	masking or weighting changes Chamfer by ≤ 0.03 cm; keeping only confidence 2 loses completeness

low-resolution indoor input, not about intrinsics-conditioned models in general.

C. Context: the price of each depth source (Exp. 1)

- 1) **The pipeline loses ~ 2 cm; LiDAR adds ~ 1 cm** (0.98 ± 0.25 cm, same in every room). The two differ clearly only at $\tau = 2$ cm: iPad LiDAR is wrong at the 2–4 cm level.
- 2) **Monocular depth with the right scale per frame (oracle) is 2.29 ± 1.11 cm behind LiDAR**, and **200 sparse metric points per frame recover it** (0.6 cm from the oracle, $3.8\times$ better than one-off calibration).
- 3) **One calibration per scene is $4.5\times$ worse than the oracle, uniformly** (19.9–26.0 cm in every room): a property of the model family, not of a failing room (§VI-B).
- 4) **The pretrained metric model is worst** (54.3 cm): the surfaces are all there, in the wrong place. Kitchen 45261556 is among its best shapes under the oracle (4.5) and its worst metric room (71.0): right shape, wrong scale. This is the baseline R3 improves on.

D. Secondary findings (Exp. 2–5)

Table VIII summarises the rest. A $13\times$ smaller model (DA-v2 Small) costs only 1.0 cm for a $6\times$ speed-up, whereas one scale per scene costs 17.6 cm: model size is not the bottleneck, scale is. Frame count governs coverage, not accuracy, and LiDAR confidence maps give no measurable gain.

TABLE IX
PER-FRAME SCALE OF THE PRETRAINED METRIC MODEL VS. R3 (SIX TEST ROOMS, EVERY FIFTH FRAME).

	Pretrained	R3
median(GT/pred) per room	0.69–0.85	0.93–1.21
per-frame swing within a room (p95/p5)	$1.55\times$	$1.52\times$
AbsRel, raw	0.381	0.134
AbsRel, one (S, T) per room from 10 GT frames	0.127	0.117
AbsRel, right scale on every frame	0.045	0.051
AbsRel, oracle (s, t) per frame	0.035	0.042




Fig. 4. Per-frame oracle scale divided by the per-scene scale (DA-v2 Large, all scenes) against the frame’s median GT depth: the scale a frame needs depends on how close the camera is to what it sees.

VI. DISCUSSION

A. Why fine-tuning works—and why it is not enough

We analysed every fifth frame of the six test rooms (471 frames): per frame, the GT/pred ratio and AbsRel under increasing amounts of scale help (Table IX). **Fine-tuning learns the domain’s multiplier.** The pretrained model is 15–31% too far in every room; R3 is within $\sim 20\%$ in both directions. R3 without calibration (0.134) is on par with the pretrained model given a per-room scale fitted on ground truth (0.127): what a one-off calibration provides has moved into the head’s weights, learned from LiDAR in other rooms. With the encoder frozen, shape barely changes (oracle 0.035 vs. 0.042). **The per-frame swing remains.** Within one room, frames at p5 and p95 still need factors $\sim 1.5\times$ apart, as before fine-tuning; with the right scale on every frame R3’s AbsRel would drop to 0.051. That information is not in a single image and does not come from a better teacher; it must come from metric data at run time. Fine-tuning and per-frame metric points are therefore complementary.

B. Scale follows the view, not time

Refitting (s, t) on every fifth frame of the relative model shows why one calibration per scene fails. Within one scan the central 90% of frames need scales $2.6\text{--}4.9\times$ apart; the correlation of scale with time flips sign between scenes (-0.33).

TABLE X
USE-CASE FIT: RECALL@ τ OVER 6 SCENES ($\checkmark \geq 0.9$, ~ 0.75 – 0.9 , $\times < 0.75$).

Use case (τ)	LiDAR	sparse	per_scene	R3	metric
overview (10 cm)	$\checkmark$ 0.97	$\checkmark$ 0.91	$\times$ 0.47	$\times$ 0.51	$\times$ 0.10
furniture (5 cm)	$\checkmark$ 0.91	$\sim$ 0.77	$\times$ 0.27	$\times$ 0.28	$\times$ 0.04
renovation (2 cm)	$\times$ 0.52	$\times$ 0.39	$\times$ 0.09	$\times$ 0.09	$\times$ 0.01

to $+0.30$), so this is *not cumulative drift*, whereas the correlation with the frame’s median depth is negative in every scene (-0.27 to -0.75 , Fig. 4). The worst frames are downward shots of table-tops at 0.5–0.8 m and frame-filling blank walls or cabinets: with no size cue the model guesses at “room scale,” and a table at 0.5 m lands at 1.5–2 m as a phantom surface. The metric model behaves the same way: right on average up to a per-room bias, wrong per frame (GT/pred from 0.43 to 1.35 within one room), but with almost no shift term—one scalar per frame recovers 97 % of its gap to the oracle. In 3D the same surface seen from several frames is placed at several distances and TSDF averages them into a thick slab.

Rooms that are hard for LiDAR are not those that are hard for monocular depth: the standard bedroom 47333774 is the easiest for LiDAR (2.4) but second-hardest for mono+oracle (6.2), because a small room scanned close to blank walls gives the model no cue. The mirrored bathroom is hard for every source whose scale is right (LiDAR 4.1, oracle 7.5), so a bathroom-renovation product should expect worse than the averages.

C. Implications for a product

Table X translates the results into product tiers. iPad LiDAR suffices for furniture placement but not for renovation measurement. On a device without LiDAR, neither one-off calibration nor a raw metric model serves any tier; R3 raises recall@10 cm from 0.10 to 0.51 with nothing at run time, a large step that is still short of the overview tier. With per-frame sparse points the same device reaches 6.0 cm, close to furniture placement. The order of investment the data suggest is therefore: **(1) per-frame scale from VIO points, (2) a fine-tuned model to remove the domain bias, (3) model size last**—investing in the model first returns <3 cm on a 22 cm base. Strategically, an app with Pro-device users can collect its own teacher during normal use (RGB, LiDAR depth and pose arrive together from ARKit), replacing a \$20k laser scanner; our data show no accuracy case for laser capture for this purpose. At 4 fps DA-v2 Large takes $3.6\times$ the video length on a laptop and DA-v2 Small $0.5\times$; fusion is $< 3\%$ of the time.

VII. LIMITATIONS AND FUTURE WORK

Reference and poses. The reference mesh is fused from Faro depth with the same poses as every row, so pose drift is excluded and a real phone system incurs additional error; completeness means complete relative to what Faro saw. **Scale**

of the evaluation. Mesh numbers come from six rooms (one used in development) and 2D numbers from 26, all captured with one iPad model in homes; offices, empty rooms, low light and other phone cameras are absent. **Fine-tuning.** Checkpoints were selected on only three validation rooms, whose ranking ($R3 > R2 > R1$) does not reproduce on the test rooms; R1 and R2 differ in the number of views as well as the label; and only the head of one model was trained with one recipe. **Proxies and bounds.** `per_scene` calibrates on ground truth and `sparse_points` uses clean LiDAR samples instead of triangulated VIO features, so both are upper bounds. Depth Pro was run through one implementation on 640×480 input. Times are unoptimised PyTorch on Apple Silicon and are meaningful only as ratios.

Future work, in order of impact: (1) combine R3 with per-frame sparse points, which the 2D analysis bounds at AbsRel 0.051; (2) a controlled teacher comparison on identical frames and a larger validation set; (3) training data and evaluation on more phone cameras, including Android; (4) evaluation with real ARKit/ARCore feature points and a reference captured by our own app.

VIII. CONCLUSION

We asked whether fine-tuning a metric monocular depth model with a phone’s own LiDAR makes it better than its pretrained self, measured against a laser scanner, through a room-reconstruction pipeline in which the depth source is the single variable. **It does, in every room we can measure:** fine-tuning only the head of DA-v2 Metric-Indoor on 24 rooms (R3) lowers Chamfer from 54.3 cm to 15.7 cm on all six held-out rooms and per-frame AbsRel from 0.410 to 0.130 on 20 further rooms, with no calibration or sensor at run time. **The teacher need not be a laser scanner:** Faro and LiDAR teachers are not measurably different, and knowing the camera’s intrinsics is no substitute for in-domain labels. **It is not yet a sensor:** iPad LiDAR gives 3.0 cm and a monocular model scaled per frame on 200 sparse points 6.0 cm; R3’s remaining error is a scale that swings with what each frame shows, which fine-tuning de-biases but cannot remove. The path to a usable product on devices without LiDAR is therefore a fine-tuned model combined with per-frame metric points from the device’s tracker.

REFERENCES

- [1] R. Ranftl, K. Lasinger, D. Hafner, K. Schindler, and V. Koltun, “Towards robust monocular depth estimation: Mixing datasets for zero-shot cross-dataset transfer,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 3, pp. 1623–1637, 2022.
- [2] L. Yang, B. Kang, Z. Huang, X. Xu, J. Feng, and H. Zhao, “Depth anything: Unleashing the power of large-scale unlabeled data,” in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 10 371–10 381.
- [3] L. Yang, B. Kang, Z. Huang, Z. Zhao, X. Xu, J. Feng, and H. Zhao, “Depth anything V2,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 37, 2024.
- [4] S. F. Bhat, R. Birkel, D. Wofk, P. Wonka, and M. Müller, “ZoeDepth: Zero-shot transfer by combining relative and metric depth,” *arXiv preprint arXiv:2302.12288*, 2023.

- [5] W. Yin, C. Zhang, H. Wang, S. Niklaus, L. Mai, S. Chen, and C. Shen, "Metric3D: Towards zero-shot metric 3D prediction from a single image," in *Proc. IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023, pp. 9043–9053.
- [6] L. Piccinelli, Y.-H. Yang, C. Sakaridis, M. Segu, S. Li, L. Van Gool, and F. Yu, "UniDepth: Universal monocular metric depth estimation," in *CVPR*, 2024.
- [7] A. Bochkovskii, A. Delaunoy, H. Germain, M. Santos, Y. Zhou, S. R. Richter, and V. Koltun, "Depth pro: Sharp monocular metric depth in less than a second," in *ICLR*, 2025.
- [8] H. Lin, S. Peng, J. Chen, S. Peng, J. Sun, M. Liu, H. Bao, J. Feng, X. Zhou, and B. Kang, "Prompting depth anything for 4K resolution accurate metric depth estimation," in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [9] B. Curless and M. Levoy, "A volumetric method for building complex models from range images," in *Proc. SIGGRAPH*, 1996, pp. 303–312.
- [10] R. A. Newcombe, S. Izadi, O. Hilliges, D. Molyneaux, D. Kim, A. J. Davison, P. Kohli, J. Shotton, S. Hodges, and A. Fitzgibbon, "KinectFusion: Real-time dense surface mapping and tracking," in *Proc. IEEE International Symposium on Mixed and Augmented Reality (ISMAR)*, 2011, pp. 127–136.
- [11] Q.-Y. Zhou, J. Park, and V. Koltun, "Open3D: A modern library for 3D data processing," *arXiv preprint arXiv:1801.09847*, 2018.
- [12] Z. Murez, T. van As, J. Bartolozzi, A. Sinha, V. Badrinarayanan, and A. Rabinovich, "Atlas: End-to-end 3D scene reconstruction from posed images," in *Proc. European Conference on Computer Vision (ECCV)*, 2020, pp. 414–431.
- [13] J. Sun, Y. Xie, L. Chen, X. Zhou, and H. Bao, "NeuralRecon: Real-time coherent 3D reconstruction from monocular video," in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021, pp. 15 598–15 607.
- [14] M. Sayed, J. Gibson, J. Watson, V. Prisacariu, M. Firman, and C. Godard, "SimpleRecon: 3D reconstruction without 3D convolutions," in *Proc. European Conference on Computer Vision (ECCV)*, 2022, pp. 1–19.
- [15] K. Tateno, F. Tombari, I. Laina, and N. Navab, "CNN-SLAM: Real-time dense monocular SLAM with learned depth prediction," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 6243–6252.
- [16] Z. Yu, S. Peng, M. Niemeyer, T. Sattler, and A. Geiger, "MonoSDF: Exploring monocular geometric cues for neural implicit surface reconstruction," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, 2022.
- [17] G. Baruch, Z. Chen, A. Dehghan, T. Dimry, Y. Feigin, P. Fu, T. Gebauer, B. Joffe, D. Kurz, A. Schwartz, and E. Shulman, "ARKitScenes: A diverse real-world dataset for 3D indoor scene understanding using mobile RGB-D data," in *Proc. Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021.
- [18] Apple Inc., "RoomPlan – apple developer documentation," <https://developer.apple.com/documentation/roomplan>, 2022, accessed 2026-09-17.